

PDD

Lecture 1: Introduction

Prepared by Jacek Sroka

Organization

- Assessment criteria
 - Labs:
 - Working during labs/homework (20 points in total)
 - Two big programming assignments (2 * 20 points)
 - Very important to obey the deadlines
penalties apply: -1 point per every started 12h period after the deadline (up to 14 points of penalty)
more than one week of delay: don't bother
 - Total 60 points from labs and minimum 30 points to pass
 - Exam:
 - Written exam with total of 40 points
 - First term (only for people who pass the labs) and Second term (for everyone, grade only from exam)
 - Easy if you attended classes and participate
- There are lots of possibilities to extend the assignments (e.g., into MSc thesis)
 - We are open to good ideas (please discuss them first with your lab teacher)
e.g. genome scale bioinformatics
 - Implementing and testing algorithms from papers
 - Working on research papers

Questions?

Materials

- Scientific papers
- Framework presentations on youtube / tutorials
- *Mining of Massive Datasets*
Anand Rajaraman, Jure Leskovec, Jeffrey D. Ullman
- Some Coursera / edX / Stanford online courses















What is Big data and why it is interesting?



What is big data?


“Big data is like teenage sex: everybody talks about it, nobody really knows how to do it, everyone thinks everyone else is doing it, so everyone claim they are doing it.”

--Dan Ariely, Duke University



Wikipedia: „*Big data is an all-encompassing term for any collection of data sets so large or complex that it becomes difficult to process them using **traditional** data processing applications.*”

- *Started with Internet companies Google/Facebook/Twitter/Amazon*
(Statistics from 2012 from Wikibon Blog)
 - *In 2008, Google was processing 20,000 terabytes of data (20 petabytes) a day*
 - *Facebook stores, accesses, and analyzes 30+ Petabytes of user generated data*
 - *100 terabytes of data uploaded daily to Facebook*
 - *According to Twitter’s own research in early 2012, it sees roughly 175 million tweets every day, and has more than 465 million accounts*
- *Big data was embraced by „brick and mortar”*
 - *banks, telecoms, insurance companies, ...*



Competitive advantage for brick and mortar business: Frequent itemset problem

- Analyze transaction log from store
- Goal: see what customers buy in the same transaction to organize store layout and plan promotions
 - Place snacks next to beer?
 - Offer antivirus software with new computer sales?
- Support: how frequently the itemset appears in the dataset
 $\text{support}(\{B,E\})=4/5$
- Association rule learning: $[\text{'milk'}]^\wedge[\text{'bread'}] \Rightarrow [\text{'butter'}]$
- $\text{Computer} \Rightarrow \text{antivirus_software}$ [support=2%, confidence=60%]
 - usefulness and certainty of discovered rules
 - 2% of all the transactions under analysis show that computer and antivirus software are purchased together: $\text{Support}(A \rightarrow B) = \text{Support_count}(A \cup B)$
 - 60% of the customers who purchased a computer also bought the software: $\text{Confidence}(A \rightarrow B) = \text{Support}(A \cup B) / \text{Support}(A)$


transaction ID	items
1	{A,C,D}
2	{B,C,E}
3	{A,B,C,E}
4	{B,E}
5	{A,B,C,E}

Competitive advantage?



What to consider
when processing Big data?

Cluster computing



Picture by: Google Inc.

Cluster computing

- Much less problems (experts do many things for us at scale and automatize)
 - Don't need to rent server rooms, pay for electricity, hire administrators
 - No network problems (also no „Layer 0” problem – see OSI model for Layers 1-7)
 - DDoS protection and monitoring (distributed denial of service protection)
 - Updating drivers, dealing with machine failures
(e.g. HDD breaks once in every 4 years, btw. its is more reliable when it is hot)
 - Note that other people may share the same computers!
- Examples
 - Miss Poland beauty pageant
 - Pokemon GO
 - CERN
 - Elections?
 - PRISM@NSA
- Processing data on clusters requires new frameworks
 - Take parallelism into account while designing algorithms (functional based API)
 - No need to rediscover the wheel (similar to DBMS, similar problems in many Google projects)
 - Be ready for failures (many computers, long lasting computation)
 - Deal with skew

Cluster computing

- New technologies
 - Better storage
 - Very strong computers (thousands of cores, terabytes of RAM)
- Software as a service (SaaS)
 - Even less problems, but also less control and higher price

Data science

- Data science – make decisions based on how we observed the world to work
- Sometimes better to use many simple models than one sophisticated (Watson on Jeopardy; ensemble methods use multiple learning algorithms to obtain better predictive performance than could be obtained from any of the constituent learning algorithms alone)
- Often we aggregate/summarize data (pagerank, clusterization – centroids i clusteroids) to more useful representation or we find important information (frequent itemsets)
- It will get worse
 - Rate of data production is growing faster than the processing speed
 - With time there will be less and less problems for which we can use reasonable algorithms due to the growing size of input data (polynomial or even linearithmic)
 - Often result guaranteed to be good approximation is good enough
 - Different hardware is good for different problems
 - Single supercomputer with specialized memory and network
 - Cluster of commodity computers
 - Single computer with lots (TB) of RAM
 - Memory needs to be initialized
 - Graphics cards

Don't make false conclusions


- Consider if the results makes sense and why
 - “data mining” = is the discovery of “models” for data where model relates to statistics, machine learning, algorithms, data bases
 - Pejorative at start (conclusion not supported by the data)
 - Bonferroni rule
 - Terrorists in hotels
W. Bush administration's Total Information Awareness data-collection and data mining plan of 2002
 - Torture the data, and it will confess to anything
-Ronald Coase, economist, Nobel Prize Laureate
 - The Bible Code
 - The results need interpretation (beer and diapers vs diapers and beer)
 - Correlation does not imply causation (Facebook i Princeton)
 - Watch for Weapons of Math Destruction
 - Teachers vs football
 - Hard to appeal black boxes
 - Use feedback loops
- Check for input data bias
 - Learn from ML projects mistakes
 - Failing to account for cultural bias can have consequences



General plan of the lecture

General plan of the lecture

- Technologies:
 - Spark, MLlib, Parquet
 - Stream processing (Kafka, Apache Beam)
 - ...
- Nice algorithmic ideas and how to apply them:
 - Locality Sensitive Hashing
 - Multiway joins, Minimal MapReduce algorithms
 - Distributing ML algorithms
 - ...
- What matters when distributing:
 - computation vs communication cost
 - skew
- How to pick the right tool to solve a problem
- How to avoid typical mistakes



Evolution of: Technologies, Frameworks, Tools, ...

Big data framework history

(good systems need companies that support them)

- Hadoop, Avro

- MapReduce
- HDFS
- YARN from 2.0



- Spark, Parquet

- Spark Core, SQL and Streaming
- GraphX, GraphFrames
- MLlib
- AMPLab, UC Berkley, Databricks



- Google Inc.

- Apache Beam
- BigQuery



- Snowflake

- Snowflake

- Nvidia

- DASK, RAPIDS



Before Big data

- Data warehouses: storing and querying historical business data
 - 1970's: Relational databases (w/SQL)
 - 1980's: (more data requires multiple computers) Parallel database systems based on “shared-nothing” architectures (Gamma, GRACE, Teradata)
 - “The term shared nothing architecture was coined by Michael Stonebraker (1986) to describe a multiprocessor database management system in which neither memory nor disk storage is shared among the processors.”
 - “Each node is independent and self-sufficient, and there is no single point of contention across the system”
 - Hash partitioning
 - 2000's: Netezza, Aster Data, DATAlegro, Greenplum, Vertica, ParAccel,...
 - 100M \$ to 1B acquisitions
- On-line transaction processing (OLTP)
 - Order entry, retail sales, and financial transaction systems, ...
 - Produces data for data warehouses
 - Shared-nothing also useful for scaling OLTP (1980's: Tandem's NonStop SQL)

Parallel database software stack

- Compiler takes query and produces query plan
- Datflow layer executes query plan
- Executed on storage managers on individual machines
- Upper layers orchestrate execution of lower layers
- Often built on top of some open source DBMS with added orchestration software: with node being instances of indexed DBMS and some
- **This stack is closed and expensive!**
- **Access data with SQL only!**

Big data in systems world

- Late 1990 – indexing and querying exploding content of the Web
 - DB technology was used but abandoned
 - Google, Yahoo! *et al* needed something else and built something
- Google
 - Google File System (GFS)
 - Files stored on 1000's of machines
 - Replication for fault-tolerance and availability
 - MapReduce (MR) programming model
 - Think about the data: user provides two simple functions (process one element + group)
 - “Parallel programming for dummies” - partitioning, message passing, fault tolerance
- Yahoo!, Facebook, and other cloned good ideas from Google's “Big Data” infrastructure
 - GFS -> Hadoop Distributed File System (HDFS)
 - MapReduce -> Hadoop MapReduce
 - Used for Web indexing, click stream analysis, log analysis, information extraction, some machine learning

MapReduce@Google

- Typical motivation: the best tools are created to avoid mundane tasks!
- Sanjay Ghemawat, Howard Gobioff, Shun-Tak Leung: **The Google file system**. SOSP 2003: 29-43
 - scalable distributed file system for large distributed data-intensive applications
 - provides fault tolerance
 - tuned to inexpensive commodity hardware
 - delivers high aggregate performance to a large number of clients
- Jeffrey Dean and Sanjay Ghemawat. 2004. **MapReduce: simplified data processing on large clusters**. In Proceedings of the 6th conference on Symposium on Operating Systems Design & Implementation - Volume 6 (OSDI'04), Vol. 6. USENIX Association, Berkeley, CA, USA, 10-10.
 - partitioning the input data
 - scheduling the program's execution across a set of machines
 - handling machine failures, and
 - managing the required inter-machine communication
 - Simple programming model: need to provide two functions **map** i **reduce** (think about a hammer, do not try to use it for every problem)

Tailor solution to your needs

- How to choose between: snowflake, spark, dask, open source, ...
- Pipeline vs ad hoc
 - Expressive/high level/well know API (SQL, pandas, scikit-learn, etc.) vs human optimization?
 - Hashing/Indexes?
- Computer vs cluster vs cloud vs hpc vs gpu
 - Choose best hardware and framework for your needs or framework for your hardware
- Do we need huge storage (for input, intermediate data, output)
 - store data well (binary vs textual, columnar formats, push down filters)
- Do we need to take failures into account (computation, storage)
- Do we need to scale up/down (on demand)
- Good docs/tutorials and responsive community (problems start with large deployments)
- Paid support
- Integration with other tools
- Distributing is not always good
 - Try single node first (local memory, no network, no framework overhead, no accidental complexity)
 - for some problems it is not easy/possible to design algorithms that scale well
 - for other it may work worse if you do not understand what is going on well (e.g. graph algorithms, wrong configuration of memory, checkpointing, etc.)

MapReduce

MapReduce

- Problems
 - Distributed programming is difficult (it's nice to have a simple model)
 - Network transfer is the main limiting factor (for 1Gbps transfer of 10TB takes 1 day)
 - Long running computations need to be resistant to single node failures
 - Need algorithms with low complexity and small data exchanged
- MapReduce
 - Simple programming model, well tested framework
 - Lots of nice algorithmic ideas available
 - Easy to use as intermediate layer for other high level languages like SQL
 - High **accidental complexity**, but lots of projects that help with that like Crunch, Cascading or Oozie
 - Data replication (based on HDFS)
 - chunk servers are compute servers
 - When possible schedule processing to a node with the data
 - Very pessimistic approach to failures
 - Big overhead for iterative algorithms (data mining, graph processing)
- We won't program in MapReduce but we will use it as introduction to Massively Parallel Computation (MPC) model
 - Goal for the first lab solve simple problems in MapReduce and implement in Spark (which has better API but don't use it yet)

Hello World

- Goal: count number of occurrences of a given word in a large text file
 - Web sever log analysis – finding popular URLs
 - Finding stop words in a given language
- What we need to take into account
 - File is too large to fit in memory
 - Pairs `<word, count>` do not fit in memory

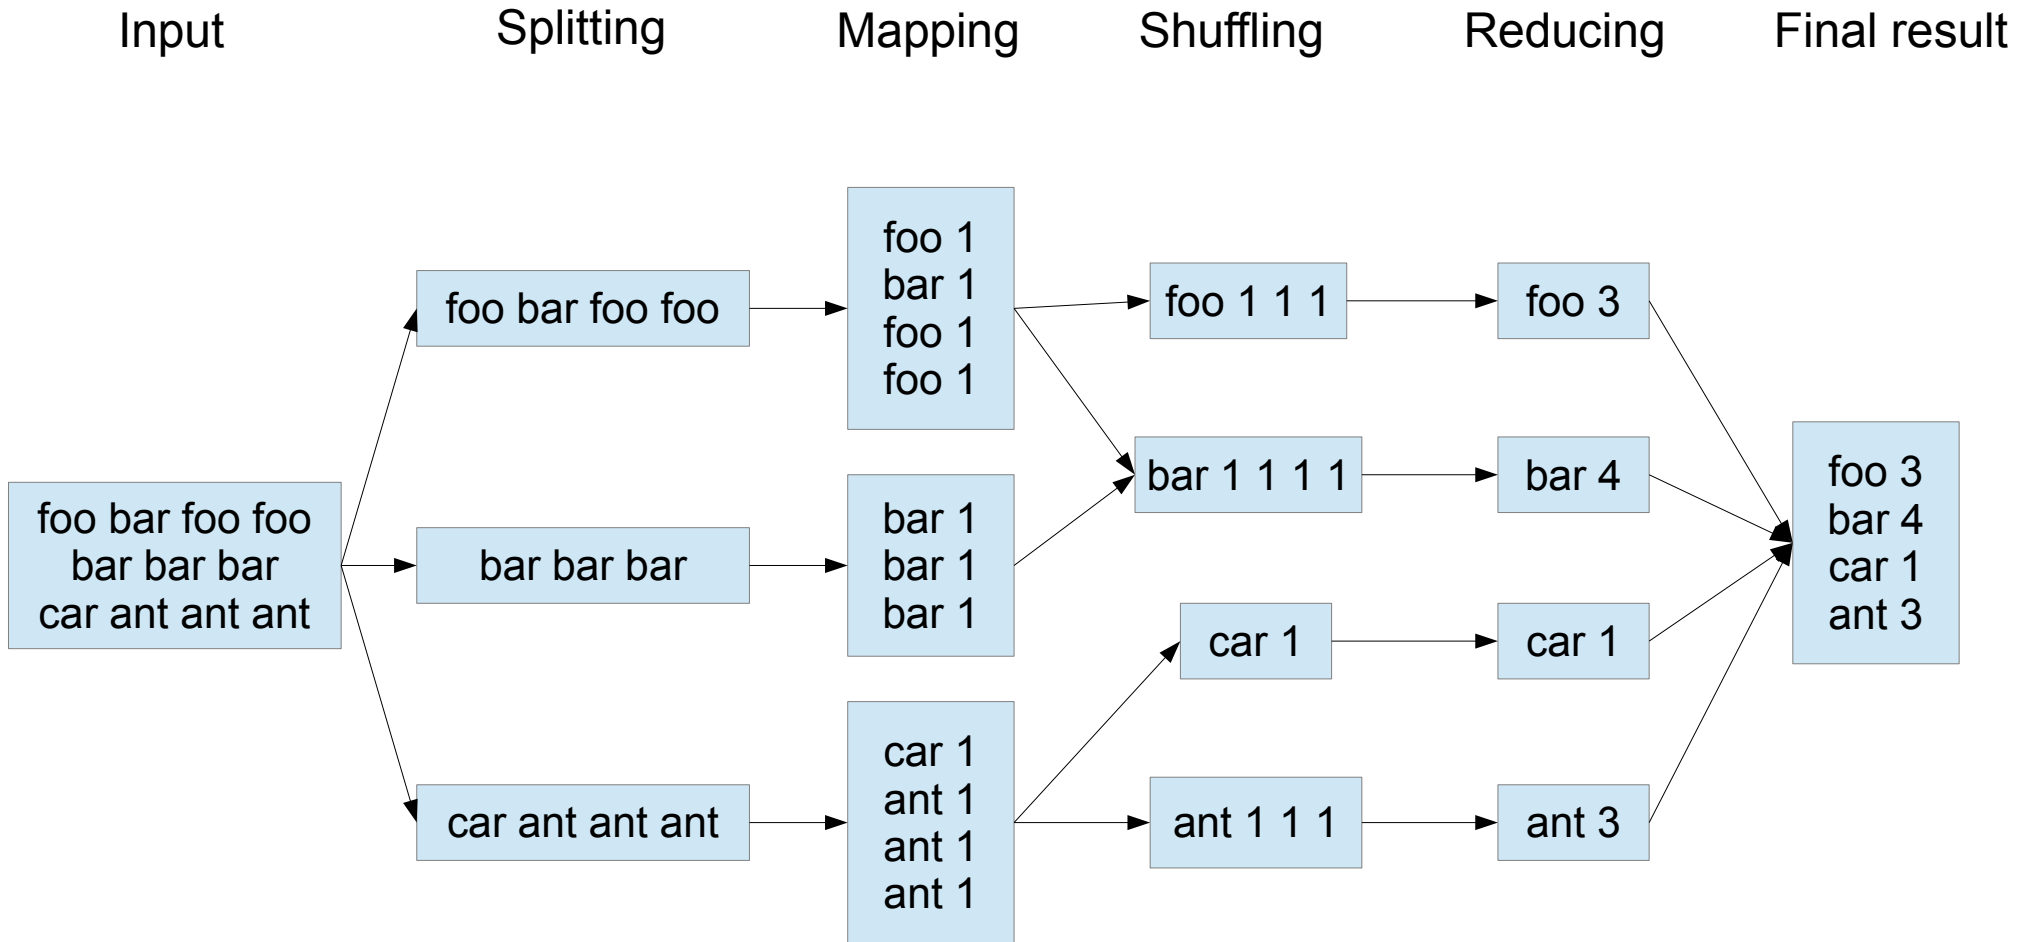
Hello World

- Goal: count number of occurrences of a given word in a large text file
 - Web sever log analysis – finding popular URLs
 - Finding stop words in a given language
- What we need to take into account
 - File is too large to fit in memory
 - Pairs `<word, count>` do not fit in memory
- Ideas for a solution
 - Use HashTable
 - Easy to distribute (process independently, then integrate results)
 - Start with some grouping or sorting, e.g., with some file system tool, then process smaller files
 - Also easy to distribute

MapReduce

- Map
 - Scan input record by record, do some preprocessing
 - Choose group for the record (based on some key) write in local file
 - Easy to distribute
- Shuffle/Sort
 - Shuffle local files so that records with the same key end up at the same node
- Reduce
 - Assuming that groups can be processed on nodes, i.e., are not materialized in memory or are small enough
- Lots of reasonable ways to choose a key for grouping (depends on the use case)
 - Whole word
 - First letter
 - 0 or 1 for the first or second part of alphabet
 - The goal is to balance the groups as much as possible

MapReduce example



MapReduce c.d.

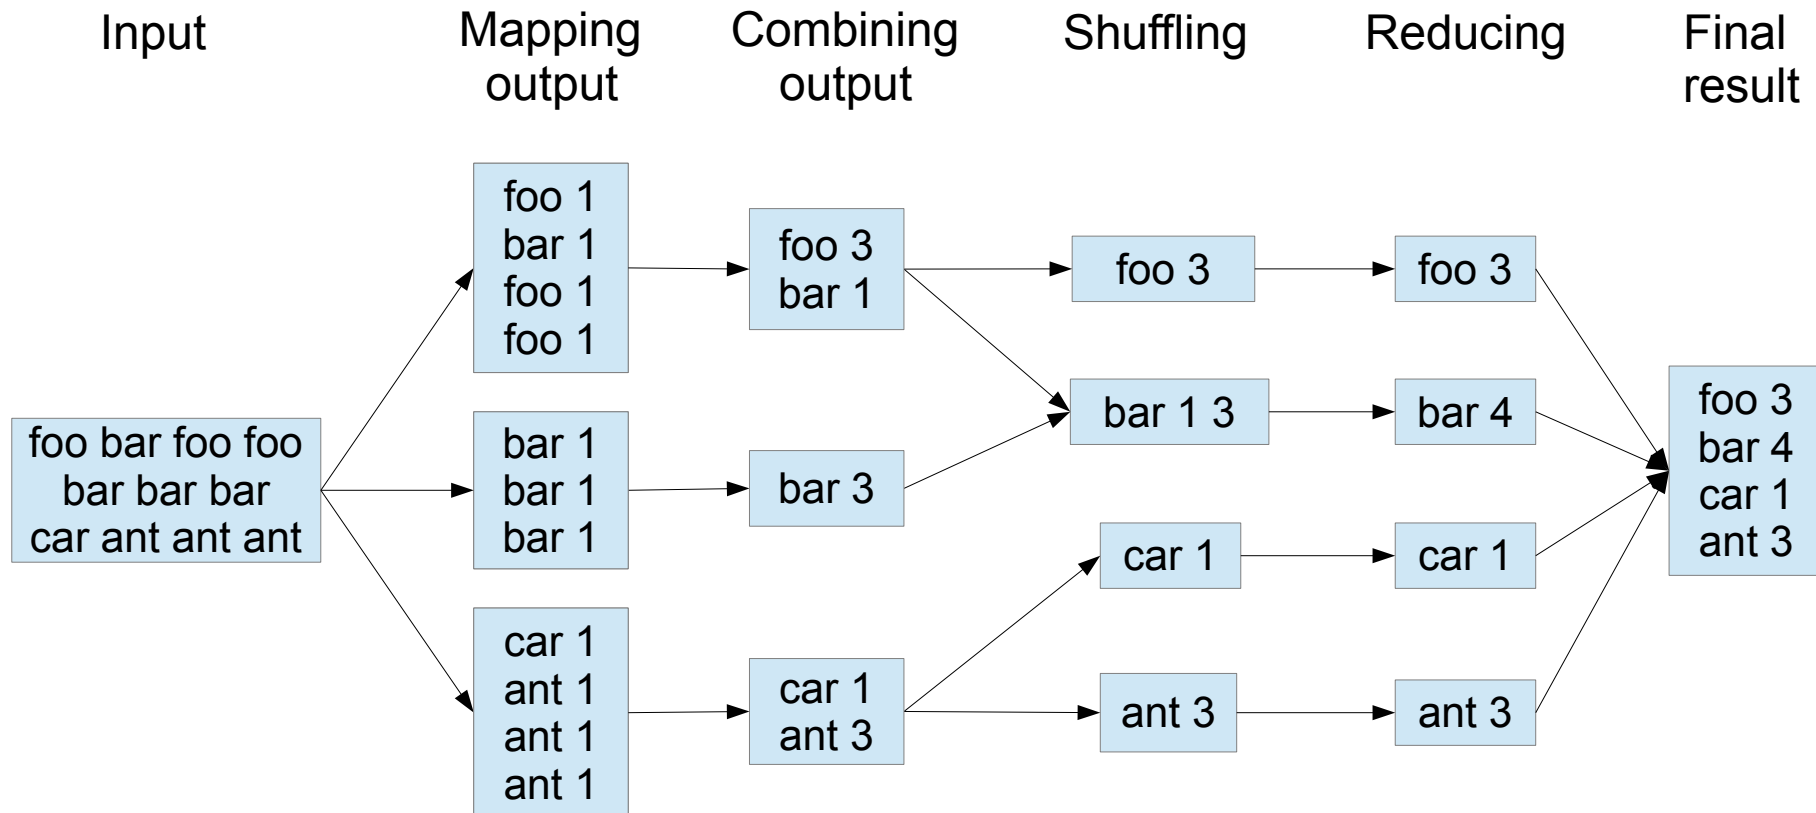
- Input: collection of key/value pairs
- $\text{Map}(k,v) \rightarrow \langle k',v' \rangle^*$
 - Execute once for each input pair
 - Results in an arbitrary number of intermediate key/value pairs
- $\text{Reduce}(k', (v1', v2', \dots)) \rightarrow \langle k'',v'' \rangle^*$
 - Execute once for each k'
 - Has access to all values for that key
- Example:
 - Read input in parts, e.g., line by line
 - map: for each line generate pairs $\langle \text{word}, 1 \rangle$
 - reduce: return length of the values list
 - map and reduce are executed on multiple nodes
- Disc access is sequential not random (VERY important for HDDs, but for SSDs too)

Terminology

- Map-Reduce job =
 - Map function (inputs -> key-value pairs)
 - +
 - Reduce function (key and list of values -> outputs)
- Map and Reduce Tasks/Processes apply Map or Reduce function to (typically) many of their inputs.
 - Unit of parallelism
- **Mapper** = application of the Map function to a single input
- **Reducer** = application of the Reduce function to a single key-(list of values) pair

Combiners

- Usually we can pre-aggregate Map task outputs with reduce function



- Works only for commutative ($a + b = b + a$) and associative ($a + (b + c) = (a + b) + c$) functions
 - e.g. for sum we can, but not for avg $(1+2+3)/3 \neq ((1+2)/2 + 3/1)/2$
 - often this can be fixed – for avg by emitting (sum, count) and using different Comb. and Red.
 - But not always – median